

《分布式编程模型与系统》期末考查作业

学号:10195501409

姓名: 卞思頔

成绩:

1. 实验目的

验证 Spark 和 MapReduce 执行 PageRank 迭代应用的性能差异

2. 设计的思路

- MapReduce 执行迭代计算过程中会反复读写 HDFS, 因此可以在 HDFS 中观察到每一轮迭代的输出结果
- MapReduce 会提交一系列的作业, 而 spark 仅有一个应用, 在 Yarn 的 UI 显示会不一样
- 对于同样规模的数据集, spark 执行时间应当更短

3. 实验设置

操作系统版本: Centos7, Hadoop3.0.0、Spark3.0.0

数据集名称: input.txt

4. 实验过程

- 针对 MapReducer 我们实现一个 PageRank, 并在 hadoop 集群上进行运行, 进行 10 轮的迭代, 查看运行的结果

```
1 package cs.author;
2
3 import java.io.IOException;
4 import java.util.StringTokenizer;
5 import org.apache.hadoop.conf.Configuration;
6 import org.apache.hadoop.fs.Path;
7 import org.apache.hadoop.io.Text;
8 import org.apache.hadoop.mapreduce.Job;
9 import org.apache.hadoop.mapreduce.Mapper;
10 import org.apache.hadoop.mapreduce.Reducer;
11 import org.apache.hadoop.mapreduce.Reducer.Context;
12 import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
13 import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
14
15 public class PageRank {
16
17     /*map过程*/
18     public static class LxnMapper extends Mapper<Object,Text,Text,Text>{
19         private String id;
20         private float pr;
21         private int count;
22         private float average_pr;
23         public void map(Object key,Text value,Context context)
24             throws IOException,InterruptedException{
25             StringTokenizer str = new StringTokenizer(value.toString()); //对value进行解析
26             id =str.nextToken(); //id为解析的第一个词, 代表当前网页
27             pr = Float.parseFloat(str.nextToken()); //pr为解析的第二个词, 转换为Float类型, 代表PageRank值
28             count = str.countTokens(); //count为剩余词的个数, 代表当前网页的出链网页个数
29             average_pr = pr/count; //求出当前网页对出链网页的贡献值
30             String linkids = "&"; //下面是输出的两类, 分别有'@'和'&'区分
```

```

34         linkids += " " + linkid;
35     }
36     context.write(new Text(id), new Text(linkids)); //输出的是<当前网页, 所有出链网页>
37 }
38 }
39
40 /*reduce过程*/
41 public static class Lxnreduce extends Reducer<Text,Text,Text,Text>{
42     public void reduce(Text key,Iterable<Text> values,Context context)
43         throws IOException,InterruptedException{
44         String lianjie = "";
45         float pr = 0;
46         /*对values中的每一个val进行分析, 通过其第一个字符是'@' 还是'&' 进行判断
47         通过这个循环, 可以 求出当前网页获得的贡献值之和, 也即是新的PageRank值; 同时求出当前
48         网页的所有出链网页 */
49         for(Text val:values){
50             if(val.toString().substring(0,1).equals("@")){
51                 pr += Float.parseFloat(val.toString().substring(1));
52             }
53             else if(val.toString().substring(0,1).equals("&")){
54                 lianjie += val.toString().substring(1);
55             }
56         }
57
58         pr = 0.8f*pr + 0.2f*0.25f; //加入跳转因子, 进行平滑处理
59         String result = pr+lianjie;
60         context.write(key, new Text(result));
61     }
62 }
63
64
65 pr = 0.8f*pr + 0.2f*0.25f; //加入跳转因子, 进行平滑处理
66 String result = pr+lianjie;
67 context.write(key, new Text(result));
68 }
69 }
70
71 @
72 public static void main(String[] args) throws Exception{
73     Configuration conf = new Configuration();
74     String pathIn1 = args[0];
75     String pathOut=args[1];
76     for(int i=1;i<41;i++){ //加入for循环
77         Job job = new Job(conf, "page rank");
78         job.setJarByClass(PageRank.class);
79         job.setMapperClass(Lxnmapper.class);
80         job.setReducerClass(Lxnreduce.class);
81         job.setOutputKeyClass(Text.class);
82         job.setOutputValueClass(Text.class);
83         FileInputFormat.addInputPath(job, new Path(pathIn1));
84         FileOutputFormat.setOutputPath(job, new Path(pathOut));
85         pathIn1 = pathOut; //把输出的地址改成下一次迭代的输入地址
86         pathOut = pathOut+i; //把下一次的输出设置成一个新地址。
87         job.waitForCompletion(true); //把System.exit()去掉
88     }
89 }
90
91 }
92
93 }

```

```

[bd@hadoop001 ~]$ hadoop jar PageRank.jar cs.author.PageRank /output/input.txt /output/output/
2022-06-23 08:46:29,719 INFO client.RMProxy: Connecting to ResourceManager at hadoop002/10.110.8.202:8032
2022-06-23 08:46:31,074 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool
with ToolRunner to remedy this.
2022-06-23 08:46:31,125 INFO mapreduce.JobResourceUploader: Disabling Erasure Coding for path: /tmp/hadoop-yarn/staging/bd/.stagi
2022-06-23 08:46:31,683 INFO input.FileInputFormat: Total input files to process : 1
2022-06-23 08:46:31,740 INFO lzo.GPLNativeCodeLoader: Loaded native gpl library from the embedded binaries
2022-06-23 08:46:31,746 INFO lzo.LzoCodec: Successfully loaded & initialized native-lzo library [hadoop-lzo rev 5dbdddb8c9b544e58
2022-06-23 08:46:32,223 INFO mapreduce.JobSubmitter: number of splits:1
2022-06-23 08:46:32,637 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1652841517163_0079
2022-06-23 08:46:32,641 INFO mapreduce.JobSubmitter: Executing with tokens: []
2022-06-23 08:46:33,153 INFO conf.Configuration: resource-types.xml not found
2022-06-23 08:46:33,154 INFO resource.ResourceUtils: Unable to find 'resource-types.xml'.
2022-06-23 08:46:33,336 INFO impl.YarnClientImpl: Submitted application application_1652841517163_0079
2022-06-23 08:46:33,460 INFO mapreduce.Job: The url to track the job: http://hadoop002:8088/proxy/application_1652841517163_0079/
2022-06-23 08:46:33,462 INFO mapreduce.Job: Running job: job_1652841517163_0079
2022-06-23 08:46:53,841 INFO mapreduce.Job: Job job_1652841517163_0079 running in uber mode : false
2022-06-23 08:46:53,846 INFO mapreduce.Job: map 0% reduce 0%
2022-06-23 08:47:08,228 INFO mapreduce.Job: map 100% reduce 0%
2022-06-23 08:47:21,413 INFO mapreduce.Job: map 100% reduce 100%
2022-06-23 08:47:22,451 INFO mapreduce.Job: Job job_1652841517163_0079 completed successfully
2022-06-23 08:47:22,726 INFO mapreduce.Job: Counters: 53
File System Counters

```

```

Shuffled Maps =1
Failed Shuffles=0
Merged Map outputs=1
GC time elapsed (ms)=218
CPU time spent (ms)=4350
Physical memory (bytes) snapshot=568655872
Virtual memory (bytes) snapshot=5197975552
Total committed heap usage (bytes)=605028352
Peak Map Physical memory (bytes)=324800512
Peak Map Virtual memory (bytes)=2597781504
Peak Reduce Physical memory (bytes)=243855360
Peak Reduce Virtual memory (bytes)=2600194048

Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

File Input Format Counters
  Bytes Read=67
File Output Format Counters
  Bytes Written=67
2022-06-23 08:54:13,180 INFO client.RMPProxy: Connecting to ResourceManager at hadoop002/10.110.8.202:8032
2022-06-23 08:54:13,200 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool
interface and execute your application with ToolRunner to remedy this.
2022-06-23 08:54:13,206 INFO mapreduce.JobResourceUploader: Disabling Erasure Coding for path: /tmp/hadoop-yarn/staging/bd/.stagi
ng/job_1652841517163_0092
2022-06-23 08:54:13,653 INFO input.FileInputFormat: Total input files to process : 1
2022-06-23 08:54:14,173 INFO mapreduce.JobSubmitter: number of splits:1

```

可以看到 10 次迭代 Mapreduce 花费了 480s 时间

- (2) 针对 Spark 我们实现一个 PageRank，并在 HDFS 上运行，在 Yarn 的 UI 上查看运行的时间和结果

```

32     val parts = s.split( regex = "\\s+")
33     (parts(0), parts(1))
34   }.distinct().groupByKey().cache()
35   var ranks = links.mapValues(v => 1.0)
36
37   for (i <- 1 to iters) {
38     val contribs = links.join(ranks).values.flatMap{ case (urls, rank) =>
39       val size = urls.size
40       urls.map(url => (url, rank / size))
41     }
42     ranks = contribs.reduceByKey(_ + _).mapValues(0.15 + 0.85 * _)
43   }
44
45   val output = ranks.collect()
46   output.foreach(tup => println(s"${tup._1} has rank: ${tup._2} ."))
47
48   spark.stop()
49 }

```



```
27
28     val iters = 10
29     val lines = spark.read.textFile(args(0)).rdd
30     val links = lines.map{ s =>
31         val parts = s.split( regex = "\\s+" )
32         (parts(0), parts(1))
33     }.distinct().groupByKey().cache()
34     var ranks = links.mapValues(v => 1.0)
35
36     for (i <- 1 to iters) {
37         val contribs = links.join(ranks).values.flatMap{ case (urls, rank) =>
38             val size = urls.size
39             urls.map(url => (url, rank / size))
40         }
41         ranks = contribs.reduceByKey(_ + _).mapValues(0.15 + 0.85 * _)
42     }
43
44     val output = ranks.collect()
45     output.foreach(tup => println(s"${tup._1} has rank: ${tup._2} ."))
46
47     spark.stop()
48 }
49
50
```

```
[bd@hadoop001 softwares]$ spark-submit --class cs.author.SparkPageRank --master yarn --deploy-mode cluster PageRank.jar
2022-06-23 09:49:19,131 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java
classes where applicable
2022-06-23 09:49:19,347 INFO client.RMPProxy: Connecting to ResourceManager at hadoop002/10.110.8.202:8032
2022-06-23 09:49:20,021 INFO yarn.Client: Requesting a new application from cluster with 4 NodeManagers
2022-06-23 09:49:21,757 INFO conf.Configuration: resource-types.xml not found
2022-06-23 09:49:21,759 INFO resource.ResourceUtils: Unable to find 'resource-types.xml'.
2022-06-23 09:49:21,812 INFO yarn.Client: Verifying our application has not requested more than the maximum memory capability of
the cluster (8192 MB per container)
2022-06-23 09:49:21,815 INFO yarn.Client: Will allocate AM container, with 2432 MB memory including 384 MB overhead
2022-06-23 09:49:21,817 INFO yarn.Client: Setting up container launch context for our AM
2022-06-23 09:49:21,819 INFO yarn.Client: Setting up the launch environment for our AM container
2022-06-23 09:49:21,856 INFO yarn.Client: Preparing resources for our AM container
2022-06-23 09:49:21,976 WARN yarn.Client: Neither spark.yarn.jars nor spark.yarn.archive is set, falling back to uploading librar
ies under SPARK_HOME.
2022-06-23 09:49:29,013 INFO yarn.Client: Uploading resource file:/tmp/spark-4f8e6b71-5600-46d5-836c-ffdacb4131ec/_spark_libs_8
470964597418922572.zip -> hdfs://hadoop001:8020/user/bd/.sparkStaging/application_1652841517163_0093/_spark_libs_84709645974189
22572.zip
2022-06-23 09:49:31,547 INFO yarn.Client: Uploading resource file:/opt/softwares/PageRank.jar -> hdfs://hadoop001:8020/user/bd/.s
parkStaging/application_1652841517163_0093/PageRank.jar
2022-06-23 09:49:32,125 INFO yarn.Client: Uploading resource file:/tmp/spark-4f8e6b71-5600-46d5-836c-ffdacb4131ec/_spark_conf_5
494815540509351277.zip -> hdfs://hadoop001:8020/user/bd/.sparkStaging/application_1652841517163_0093/_spark_conf_.zip
2022-06-23 09:49:32,270 INFO spark.SecurityManager: Changing view acls to: bd
2022-06-23 09:49:32,272 INFO spark.SecurityManager: Changing modify acls to: bd
2022-06-23 09:49:32,274 INFO spark.SecurityManager: Changing view acls groups to:
2022-06-23 09:49:32,276 INFO spark.SecurityManager: Changing modify acls groups to:
2022-06-23 09:49:32,277 INFO spark.SecurityManager: SecurityManager: authentication disabled; ui acls disabled; users with view
permissions: Set(bd); groups with view permissions: Set(); users with modify permissions: Set(bd); groups with modify permission
s: Set()
2022-06-23 09:49:32,492 INFO yarn.Client: Submitting application application_1652841517163_0093 to ResourceManager
```

```
+ hadoop201
2022-06-23 09:50:06,820 INFO yarn.Client: Application report for application_1652841517163_0093 (state: RUNNING)
2022-06-23 09:50:07,824 INFO yarn.Client: Application report for application_1652841517163_0093 (state: RUNNING)
2022-06-23 09:50:08,828 INFO yarn.Client: Application report for application_1652841517163_0093 (state: RUNNING)
2022-06-23 09:50:09,832 INFO yarn.Client: Application report for application_1652841517163_0093 (state: RUNNING)
2022-06-23 09:50:10,837 INFO yarn.Client: Application report for application_1652841517163_0093 (state: RUNNING)
2022-06-23 09:50:11,841 INFO yarn.Client: Application report for application_1652841517163_0093 (state: RUNNING)
2022-06-23 09:50:12,845 INFO yarn.Client: Application report for application_1652841517163_0093 (state: RUNNING)
2022-06-23 09:50:13,849 INFO yarn.Client: Application report for application_1652841517163_0093 (state: RUNNING)
2022-06-23 09:50:14,853 INFO yarn.Client: Application report for application_1652841517163_0093 (state: RUNNING)
2022-06-23 09:50:15,857 INFO yarn.Client: Application report for application_1652841517163_0093 (state: RUNNING)
2022-06-23 09:50:16,862 INFO yarn.Client: Application report for application_1652841517163_0093 (state: RUNNING)
2022-06-23 09:50:17,866 INFO yarn.Client: Application report for application_1652841517163_0093 (state: RUNNING)
2022-06-23 09:50:18,869 INFO yarn.Client: Application report for application_1652841517163_0093 (state: RUNNING)
2022-06-23 09:50:19,878 INFO yarn.Client: Application report for application_1652841517163_0093 (state: RUNNING)
2022-06-23 09:50:19,879 INFO yarn.Client:
client token: N/A
diagnostics: AM container is launched, waiting for AM container to Register with RM
ApplicationMaster host: N/A
ApplicationMaster RPC port: -1
queue: default
start time: 1655948972534
final status: UNDEFINED
tracking URL: http://hadoop002:8088/proxy/application_1652841517163_0093/
user: bd
2022-06-23 09:50:20,885 INFO yarn.Client: Application report for application_1652841517163_0093 (state: ACCEPTED)
2022-06-23 09:50:21,890 INFO yarn.Client: Application report for application_1652841517163_0093 (state: ACCEPTED)
2022-06-23 09:50:22,895 INFO yarn.Client: Application report for application_1652841517163_0093 (state: ACCEPTED)
2022-06-23 09:50:23,906 INFO yarn.Client: Application report for application_1652841517163_0093 (state: ACCEPTED)
2022-06-23 09:50:24,923 INFO yarn.Client: Application report for application_1652841517163_0093 (state: ACCEPTED)
2022-06-23 09:50:25,929 INFO yarn.Client: Application report for application_1652841517163_0093 (state: ACCEPTED)
```

A	0.15	B C D
B	0.21666667	A D
C	0.41666667	C
D	0.21666667	B C

可以发现 spark 执行了 100s 左右，

(3) 查看 MapReducer 和 Spark 运行的时间的差异，得出结论。

5. 结论

在运行中我们可以发现，在迭代 10 次的情况下，Spark 的运行时间为 Mapreduce 的仅为 1/4 到 1/5，如果增加迭代次数或者增加数据量，则运行时间能更加大大缩减，spark 能对 mapreduce 存在极大的性能优势。

mapreduce 是基于磁盘进行计算，与磁盘存在大量的磁盘 IO 交换，而 Spark 是基于内存计算，基于 DAG 的任务调度机制，则能大大提升计算的速度和效率。

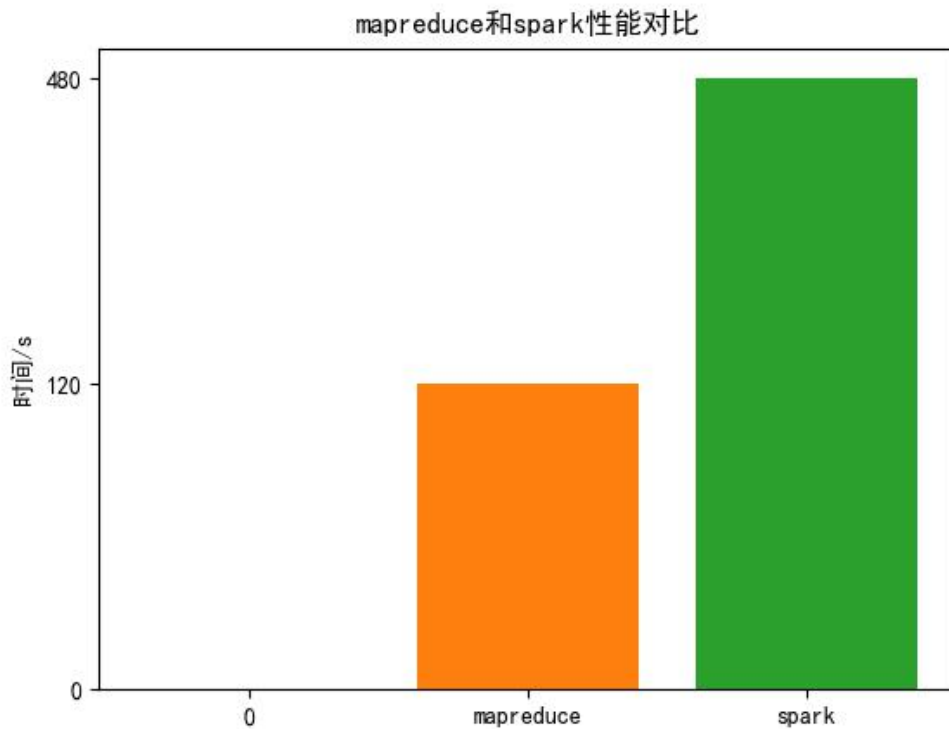


图 1 mapreduce 和 spark 进行 pagerank 运算时间对比

MapReduce	Spark
数据存储结构：磁盘HDFS文件系统的split	使用内存构建弹性分布式数据集RDD 对数据进行运算和cache
编程范式：Map + Reduce	DAG: Transformation + Action
计算中间结果落到磁盘，IO及序列化、反序列化代价大	计算中间结果在内存中维护 存取速度比磁盘高几个数量级
Task以进程的方式维护，需要数秒时间才能启动任务	Task以线程的方式维护 对于小数据集读取能够达到亚秒级的延迟

图 2 mapreduce 和 spark 进行运算时的区别

附录：

实验代码地址：

<https://github.com/Emma-0129/Distributed-System-Experiment>